

Sowmith Sripadi

+1 (786)-784-2900 | Sripadi.sowmith01@gmail.com | linkedin.com/in/sripadi

PROFESSIONAL SUMMARY

Results-driven Data Engineer with almost 5 years of experience designing and delivering scalable data pipelines, cloud-native ETL workflows, and enterprise reporting platforms across banking, HR, and retail domains. Proven expertise in Azure (ADF, Synapse, ADLS Gen2), AWS (Glue, Redshift, Kinesis), and big data technologies (PySpark, Spark, Delta Lake). Adept at partnering with cross-functional stakeholders to translate complex business requirements into production-grade data solutions. Holds an M.S. in Computer Science and a Microsoft Azure Fundamentals certification.

PROFESSIONAL EXPERIENCE

Bank of America

Feb 2025 - Present

Data Engineer | Tech: Azure Data Factory | Synapse Analytics | ADLS Gen2 | PySpark | SQL | Power BI | Azure Key Vault

Project: Enterprise Credit Risk Reporting Platform

- Architected end-to-end Azure Data Factory pipelines to consolidate transaction, repayment, delinquency, and customer master data from 10+ upstream systems into a centralized Azure Data Lake Gen2 environment, reducing data onboarding time by 60%.
- Built scalable PySpark and SQL transformation logic in Azure Synapse to generate curated risk-reporting datasets, automating daily delivery and eliminating hours of manual data preparation.
- Designed and maintained Synapse warehouse tables for key risk metrics (e.g., utilization ratio, days past due, payment trends, and portfolio segmentation), enabling consistent, trusted reporting across risk and operations teams.
- Optimized ETL performance through partition-based processing, incremental loads, and query tuning, delivering portfolio insights and improving SLA adherence by 35%.
- Developed Power BI dashboards adopted by risk leadership to track portfolio health, delinquency movement, and early warning indicators for high-risk customer segments, driving faster decision cycles.
- Enforced enterprise security controls including PII masking, RBAC, and Azure Key Vault secret management, ensuring compliance with SOX and internal banking data governance standards.
- Partnered with risk, compliance, and analytics teams to validate reporting outputs and align business definitions, reducing data discrepancy tickets by 40%.

Accenture

Aug 2021 - Dec 2022

Data Engineer | Tech: Databricks | PySpark | SQL | Azure Data Factory | Azure SQL | Delta Lake | Power BI

Project: HR Analytics

- Engineered ADF pipelines aggregating HR data from 25+ global regions, reducing reporting turnaround from three days to under six hours through scheduled refresh and retry mechanisms.
- Tuned PySpark jobs in Databricks by pruning partitions and implementing intelligent caching strategies, reducing notebook runtime by 67% and accelerating transformation of employee records and policy data.
- Optimized Azure SQL queries using filtered indexes and execution plan analysis, cutting dashboard response time from 20+ seconds to under 7 seconds a 65% improvement in end-user experience.
- Automated daily incremental Power BI dataset refreshes via ADF integration, eliminating manual intervention and boosting report accuracy across compliance dashboards.
- Implemented data masking and AES-256 encryption on PII fields in alignment with GDPR and HIPAA, protecting sensitive employee health and financial data across all pipelines.

Project: Enterprise Reporting System

- Consolidated Oracle, SAP, and Snowflake data into unified, standardized datasets, enabling enterprise-wide reporting for 25,000+ employees with consistent KPI definitions across all departments.
- Built a PySpark-based data comparison framework to validate source-to-target migrations, reducing QA effort by 70% and achieving 100% reconciliation accuracy during UAT.
- Enhanced Delta Lake performance through small file compaction and partition strategy redesign, improving query response speed by 40% for business-critical dashboards.

- Designed and delivered five role-specific Power BI dashboards with secure export controls, doubling executive and analyst usage within one month of deployment.
- Led end-to-end project delivery across design, development, and implementation phases, achieving 99% production accuracy with zero rollbacks and high stakeholder satisfaction.

Tata Consultancy Services (TCS)

Jul 2019 - Aug 2021

Data Engineer | Tech: AWS Glue | SparkSQL | Amazon Redshift | AWS Kinesis | Terraform | Apache Airflow

- Processed 100M+ records daily using AWS Glue, SparkSQL, and Redshift, improving data availability and reducing pipeline latency by 40% to consistently meet SLA requirements.
- Integrated real-time AWS Kinesis streams with EMR clusters, enabling sub-minute event monitoring and reducing manual analyst review time by 30%, directly improving fraud detection capabilities.
- Developed reusable Terraform IaC templates for AWS Lambda, IAM, and S3, cutting environment provisioning time by 50% and eliminating manual deployment errors across dev and production.
- Maintained and monitored 20+ Apache Airflow DAGs with automated retries, logging, and alerting, driving measurable pipeline stability and ensuring on-time data delivery for downstream analytics teams.

EDUCATION

Master of Science, Computer Science

Jan 2023 - Dec 2024

Florida International University, Miami, FL

CERTIFICATIONS

Microsoft Certified: Azure Fundamentals (AZ-900)

TECHNICAL SKILLS

Programming & Query Languages: Python | SQL | Shell Scripting | SparkSQL | Scala (basic)

Azure Data Stack: Azure Data Factory | Synapse Analytics | ADLS Gen2 | Azure Blob Storage | Azure Key Vault | Microsoft Fabric | Dataverse | Power BI | Power Automate

ETL & Data Engineering: Apache Airflow | ADF Pipelines | Synapse Pipelines | dbt | Informatica | SSIS | Azure Stream Analytics | Event Hub

Big Data & Distributed Computing: Apache Spark (PySpark) | Hadoop | Hive | Kafka | Delta Lake | Microsoft Fabric Lakehouse

Databases: Azure SQL Database | OracleDB | MongoDB | MySQL | PostgreSQL | DynamoDB

Data Warehousing: Snowflake | Azure Synapse | Google BigQuery | Amazon Redshift

Data Modeling & File Formats: Star/Snowflake Schema | Dimensional Modeling | ER Modeling | Data Vault | Parquet | Avro | JSON | CSV

Data Governance & Security: RBAC | IAM | VPC | Azure Purview | KMS | PII Masking | AES-256 Encryption | GDPR | HIPAA | SOX

Visualization & Reporting: Power BI | Tableau | QuickSight | Azure Monitor | Log Analytics

DevOps & CI/CD: Git | GitHub | GitHub Actions | Docker | Terraform | Azure DevOps Pipelines

Methodologies: Agile/Scrum | Stakeholder Communication | Cross-functional Collaboration